

YANG CHEN

NZ: (+64) 022-5250016 / CN: (+86) 185-1182-2753

✉ yang.chen@auckland.ac.nz • 🌐 www.yangchen.info

EDUCATION

University of Auckland

Ph.D. in Computer Science

Auckland, New Zealand

November 2018 – September 2022

- Supervisors: Jiamou Liu and Bakhadyr Khoussainov
- Thesis: *From One to Infinity: New Algorithms for Reinforcement Learning and Inverse Reinforcement Learning*

University of Auckland

First Class Honours in Computer Science

Auckland, New Zealand

July 2017 – July 2018

- Dissertation: *Network Building: Methodological Foundations and Algorithmic Analysis*

Beijing Institute of Technology

Bsc in Computer Science & Technology

Beijing, China

August 2013 – June 2017

WORK EXPERIENCE

University of New South Wales

Senior Research Associate

Sydney, Australia

To Start in October 2024

University of Auckland

Research Fellow

Auckland, New Zealand

September 2022 – Present

DAMO Academy, Alibaba

Research Intern

Beijing, China

September 2020 – January 2021

RESEARCH INTERESTS

My research interests span a spectrum of reinforcement learning, multi-agent systems and game theory. Recently, I have aimed to solve issues modelled as multi-agent systems from reinforcement learning and game-theoretic perspectives. Along this line, my ultimate goal is to reinforce the insights into theories of applying reinforcement learning in conjunction with game theory. An important thread of my research is reinforcement learning with many agents, where I attempt to explore exciting methods in such scenarios by combining reinforcement learning and the mean-field theory. Lately, my focus has moved to applying reinforcement learning and large language models in reasoning, behaviour modelling, and, more broadly, ethical AI and responsible AI.

TEACHING

Postgraduate Courses

- **COMPSCI 713: Artificial Intelligence Foundations**

Instructor, University of Auckland.

Semester 1, 2024

- **COMPSCI 761: Advanced Topics in Artificial Intelligence**
Instructor, University of Auckland. *Semester 2, 2022 – 2023*
- **COMPSCI 769: Natural Language Processing**
Guest Lecturer, University of Auckland. *Semester 2, 2024*

Undergraduate Courses

- **COMPSCI 367: Artificial Intelligence**
Guest Lecturer, University of Auckland. *Semester 2, 2024*
- **COMPSCI 220: Algorithms and Data Structures**
Guest Lecturer, University of Auckland. *Semester 1, 2022*

SUPERVISION

One BSc (Honours) Student (graduated); four master students; one PhD Student (mentoring).

GRANTS

Current Research

- **AI-based behavioural analytics for live sports broadcast** *2024 – 2026*
Associate Investigator
Principal Investigator: Patrice Delmas
Funding Agency: Ministry of Business, Innovation & Employment (MBIE) Endeavour Fund
Amount Awarded: 1 Million NZD

EVENT ORGANISING

- Local co-chair of AAMAS 2024.
- Co-chair of AAMAS 2024 AAAI track.
- Webmaster of AAMAS 2022, BSCI 2020, 2021, 2022, 2023.

ACADEMIC SERVICES

- **Conference Reviewer:** ACL 2024, EMNLP 2024, ECAI 2024, AAMAS 2023 – 2025.
- **Journal Reviewer:** JMLR.

AWARDS

- AAMAS 2022 Scholarship *April 2022*
- Google Global PhD Fellowship Nomination (Australia & New Zealand) *August 2020*
- Best Paper Award, BSCI 2019. *July 2019*
- Summer Scholarship Funding from PDH Research Partnership. *November 2018*
- University of Auckland Doctoral Scholarship. *October 2018*

SELECTED PUBLICATIONS

- ***(Multi-agent) Reinforcement Learning & Imitation Learning***
- **Multi-Agent, Human-Agent and Beyond: A Survey on Cooperation in Social Dilemmas**
Hao Guo, Chunjiang Mu, **Yang Chen**, Chen Shen, Shuyue Hu, Zhen Wang. *Neurocomputing Journal*. 2024.
- **Meta-Inverse Reinforcement Learning for Mean Field Games with Probabilistic Context Variables**
Yang Chen, Xiao Lin, Bo Yan, Libo Zhang, Jiamou Liu, Neset Özkan Tan, Michael Witbrock. *Thirty-Eighth AAAI Conference on Artificial Intelligence*. **AAAI 2024**. (Core A*)
- **Multi-Agent, Human-Agent and Beyond: A Survey on Cooperation in Social Dilemmas**. Hao Guo, Chunjiang Mu, **Yang Chen**, Chen Shen, Shuyue Hu, Zhen Wang. *Neurocomputing*. **2024**.
- **Adversarial Inverse Reinforcement Learning for Mean Field Games**
Yang Chen, Libo Zhang, Zhenyun Deng, Neset Özkan Tan, Jiamou Liu, Michael Witbrock. *The 22nd International Conference on Autonomous Agents and Multi-agent Systems*. **AAMAS 2023**. (Core A*)
- **Density-based Correlated Equilibrium for Markov Games**.
Libo Zhang, **Yang Chen (co-first & contact)**, Toru Takisaka, Bakh Khoussainov, Michael Witbrock, Jiamou Liu. *The 22nd International Conference on Autonomous Agents and Multi-agent Systems*. **AAMAS 2023**. (Core A*)
- **Individual-Level Inverse Reinforcement Learning for Mean Field Games**
Yang Chen, Libo Zhang, Jiamou Liu, Shuyue Hu. *The 21st International Conference on Autonomous Agents and Multi-agent Systems*. **AAMAS 2022**. (Core A*)
- **Interconnected Neural Linear Contextual Bandits with Upper Confidence Bound Exploration**
Yang Chen, Miao Xie, Jiamou Liu, Kaiqi Zhao. *26th Pacific-Asia Conference on Knowledge Discovery and Data Mining*. **PAKDD 2022**.
- **Social Structure Emergence: A Multi-agent Reinforcement Learning Framework for Relationship Building**
Yang Chen, Jiamou Liu, He Zhao, Hongyi Su. *The 19th International Conference on Autonomous Agents and Multi-agent Systems*. **AAMAS 2020**. (Core A*)
- **Social Capital Games as A Framework for Social Structural Pattern Emergence**
Yang Chen, Jiamou Liu. *IEEE/ACM International Conference on Advances in Social Networks Analysis and Mining*. **ASONAM 2020**.
- ***Multi-agent Behaviour Modelling & Simulation***
- **Behaviour Modelling of Social Animals via Causal Structure Discovery and Graph Neural Networks**

Gaël Gendron (co-first), **Yang Chen** (co-first), Mitchell Rogers, Yiping Liu, Mihailo Azhar, Shahrokh Heidari, David Arturo Soriano Valdez, Kobe Knowles, Padriac O’Leary, Simon Eyre, Michael Witbrock, Gillian Dobbie, Jiamou Liu and Patrice Delmas. *The 23rd International Conference on Autonomous Agents and Multi-agent Systems*. **AAMAS 2024**. (Core A*)

- **Meerkat Behaviour Recognition Dataset**

Mitchell Rogers, Gaël Gendron, David Soriano Valdez, Mihailo Azhar, **Yang Chen**, Shahrokh Heidari, Caleb Perelini, Padriac O’leary, Kobe Knowles, Izak Tait, Simon Eyre, Michael Witbrock, Patrice Delmas. *3rd Workshop on CV4Animals: Computer Vision for Animal Behavior Tracking and Modeling (in conjunction with CVPR 2023)*.

- **MSDC: Non-intrusive Load Monitoring with a Dual-CNN Model**

Jialing He, Jiamou Liu, Zijian Zhang, **Yang Chen**, Yiwei Liu, Bakh Khoussainov, Liehuang Zhu. *Thirty-Seventh AAAI Conference on Artificial Intelligence*. **AAAI 2023**. (Core A*)

- *Natural Language Processing*

- **Assessing and Enhancing the Robustness of Large Language Models with Task Structure Variations for Logical Reasoning**

Qiming Bao, Gaël Gendron, Alex Peng, Wanjun Zhong, Neset Tan, **Yang Chen**, Michael Witbrock, Jiamou Liu. *The 31st International Conference on Neural Information Processing*. **ICONIP 2024**.

- **Abstract Meaning Representation-Based Logic-Driven Data Augmentation for Logical Reasoning**

Qiming Bao, Alex Yuxuan Peng, Zhenyun Deng, Wanjun Zhong, Gael Gendron, Timothy Pistoletti, Neset Tan, Nathan Young, **Yang Chen**, Yonghua Zhu, Paul Denny, Michael Witbrock, Jiamou Liu. *The 62nd Annual Meeting of the Association for Computational Linguistics*. **ACL 2024 Findings**.

- **Neuromodulation Gated Transformer.**

Kobe Knowles, Joshua Bensemann, Diana Benavides Prado, Vithya Yogarajan, Michael Witbrock, Gillian Dobbie, **Yang Chen**. *The Eleventh International Conference on Learning Representations*. **ICLR 2023 Tiny Papers**.

- **Multi2Claim: Generating Scientific Claims from Multi-Choice Questions for Scientific Fact-Checking.**

Neset Tan, Trung Nguyen, Josh Bensemann, Alex Peng, Qiming Bao, **Yang Chen**, Mark Gahagan, Michael Witbrock. *The 17th Conference of the European Chapter of the Association for Computational Linguistics*. **EACL 2023**.

- **Contrastive Learning with Logic-driven Data Augmentation for Logical Reasoning over Text.**

Qiming Bao, Alex Yuxuan Peng, Zhenyun Deng, Wanjun Zhong, Neset Tan, Nathan Young, **Yang Chen**, Yonghua Zhu, Michael Witbrock, Jiamou Liu. *Symposium on Large Language Models IJCAI’23*. **LLM@IJCAI’23**.

- **Interpretable AMR-Based Question Decomposition for Multi-hop Question Answering.**

Zhenyun Deng, Yonghua Zhu, **Yang Chen**, Michael Witbrock, Patricia Riddle. *The 31st International Joint Conference on Artificial Intelligence. IJCAI 2022*. (Core A*)

- **Prompt-based Conservation Learning for Multi-hop Question Answering.**

Zhenyun Deng, Yonghua Zhu, **Yang Chen**, Qianqian Qi, Michael Witbrock, Patricia Riddle. *The 29th International Conference on Computational Linguistics. COLING 2022*. (Core A)

- **An explainability analysis of a sentiment prediction task using a transformer-based attention filter**

Neset Özkan Tan, Joshua Bensemann, Diana Benavides-Prado, **Yang Chen**, Mark Gahegan, Lia Lee, Alex Yuxuan Peng, Patricia Riddle, Michael Witbrock. *The Ninth Annual Conference on Advances in Cognitive Systems. ACS 2021*.

- **Graph Theory and Graph Neural Networks**

- **Graph Transformer against Graph Perturbation by Flexible-pass Filter**

Yonghua Zhu, Jincheng Huang, **Yang Chen**, Robert Amor, Michael Witbrock. *Journal of Information Fusion. 2024*.

- **Robust Node Classification on Graph Data with Graph and Label Noise**

Yonghua Zhu, Lei Feng, Zhenyun Deng, **Yang Chen**, Robert Amor, Michael Witbrock. *Thirty-Eighth AAAI Conference on Artificial Intelligence. AAAI 2024*. (Core A*)

- **Efficient Size-Prescribed k -Core Search**

Yiping Liu, Bo Yan, Bo Zhao, Hongyi Su, **Yang Chen**, Michael Witbrock. *The 2023 IEEE/ACM International Conference on Advances in Social Networks Analysis and Mining. ASONAM 2023*.

- **Chain of Propagation Prompting for Node Classification**

Yonghua Zhu, Zhenyun Deng, **Yang Chen**, Robert Amor, Michael Witbrock. *ACM MultiMedia 2023. ACM MM 2023*. (Core A*)

- **A Reinforcement Learning Approach to Gaining Social Capital with Partial Observation**

He Zhao, Hongyi Su, **Yang Chen**, Jiamou Liu, Hong Zheng, Bo Yan. *The 16th Pacific Rim International Conference on Artificial Intelligence. PRICAI 2019*.

- **Dynamic Relationship Building: Exploitation Versus Exploration on a Social Network**

Bo Yan, **Yang Chen**, Jiamou Liu. *The 18th International Conference on Web Information Systems Engineering. WISE 2017*.

- **Can Reinforcement Learning Enhance Social Capital?**

He Zhao, Hongyi Su, **Yang Chen**, Jiamou Liu, Bo Yan, Hong Zheng. *The International Workshop on Web Information Systems in the Era of AI. 2019*.

- **Distributed Community Detection over Blockchain Networks Based on Structural Entropy**
Yang Chen, Jiamou Liu. *The 2019 ACM International Symposium on Blockchain and Secure Critical Infrastructure*. **BSCI 2019**. **(Best Paper Award)**
- **Becoming Gatekeepers Together with Allies: Collaborative Brokerage over Social Networks**
Yang Chen, Jiamou Liu. *The 2019 IEEE/ACM International Conference on Advances in Social Networks Analysis and Mining*. **ASONAM 2019**.